The "CIFAR-110"

Image Classification Analysis via Convolutional Neural Networks

John A. Fonte, Esq.

Table of Contents

1. INTRODUCTION

- a. The Problem & Significance of Research
- b. Explanation of the Dataset

2. DATA SETUP & PREPROCESSING

- a. Data Loading
- b. Train/Test Split & Shuffling

3. DATA EXPLORATION

- a. Visualization of Images
- b. Example Image Modifications

4. DATA MODELING

- a. Experimentation with Dimensionality Reduction techniques
- b. Testing of those techniques via MLP

5. UNSUPERVISED MODELING

- a. Spectral Clustering
- b. t-SNE Clustering
- c. LDA & Comparison to t-SNE, PCA

6. SUPERVISED MODELING

- a. Simple Convolutional Neural Network
- b. Multi-Conv CNN
- c. Small Layer CNN
- d. Combination CNN

7. CONCLUSION


- a. Summary Graphs
- b. Final Analysis & Recommendations

1. INTRODUCTION

a. The Problem & Its Significance

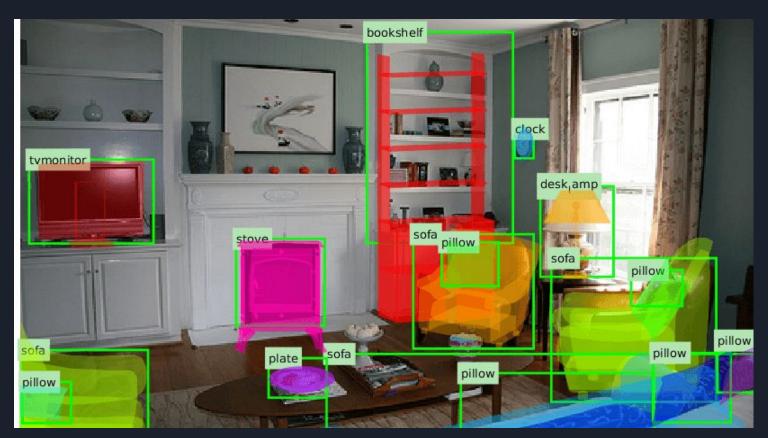
- Image Classification and Analysis has wide array of applications.
 - GPS and Geospatial Analysis
 - Facial Recognition
 - Radar and Depth Proximity (think auto-driving)





And one more widespread application...

OBJECT RECOGNITION



Problems with Object Recognition

- 1. OR Models are Becoming More Sophisticated
 - a. Identifying more kinds of objects
 - b. Identifying more objects with higher specificity
 - i. <u>Example:</u> Going from "Soda Can" to "Pepsi Can" or "Diet Coca-Cola" can
 - ii. Identifying New Classes requires a re-training of the algorithm
- 2. OR Models are Highly Complex
 - a. Almost always use Artificial Neural Networks
 - i. Training Takes Forever
 - ii. High Computational Resources are expended

Definition of Problem:



Which Object Recognition Model will Achieve the

Highest Efficiency?

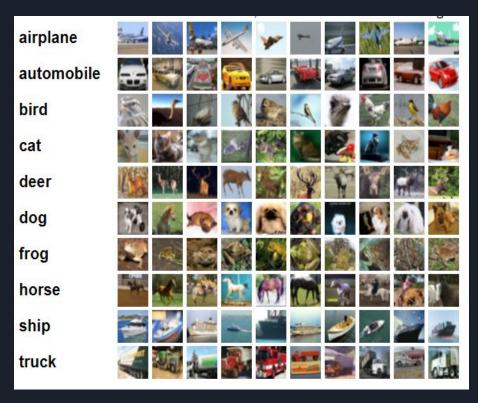
(Efficiency Defined Here: Accuracy-to-Time Rate)

Intended User: Company-side product, for companies to implement and adapt for their own user-base.

1. Introduction

b. Explanation of the Dataset: "CIFAR"

- Canadian Institute For Advanced Research
 - a. CIFAR-10
 - i. 60000 images
 - 1. 50000 train, 10000 test
 - ii. Balanced Classes
 - 1. Total of 6000 images / class
 - iii. Uniform shape & Size
 - 1. 32 x 32 x 3(RGB channels)
 - 2. Typical 0 (darkest) to 255 (brightest) metric
 - b. CIFAR-100
 - i. Same Size (60000 images)
 - ii. Balanced Classes
 - 1. 600 images / class
 - iii. Uniform Shape & Size (32 x 32 x 3)



Datasets applied here: The "CIFAR-110"

- CIFAR-110 is CIFAR-10 with addition of CIFAR-10 classes
 - Additional classes = additional complexity
 - o Switching it up from an otherwise commonly used image classification dataset
 - Both CIFAR-10 and CIFAR-110 can be loaded directly via from keras.datasets import cifar10
- Keeping CIFAR-10 for fast analysis and comparison of results to CIFAR-110

2. a. Data Loading

1. Steps:

- a. Downloaded compressed data files from CIFAR website
- b. Unzipped files
- c. Loaded files into Python's <u>Pickle</u> loading Library: and "unpickled" them
 - i. Each file yielded a dictionary, with each dictionary key being a different set of array values

```
meta10 = unpickle('D:/Github/
meta10

{b'num_cases_per_batch': 10000,
b'label_names': [b'airplane',
b'automobile',
b'bird',
b'cat',
b'deer',
b'dog',
b'frog',
b'horse',
b'ship',
b'truck'],
b'num_vis': 3072}
```

```
# checking the characteristics to import
print(batch100train.keys())
print(batch100test.keys())

dict_keys([b'filenames', b'batch_label', b'fine_labels', b'coarse_labels', b'data'])
dict_keys([b'filenames', b'batch_label', b'fine_labels', b'coarse_labels', b'data'])
```

2. Data Transformation

- a. Image Arrays Loaded as 1-D objects
- b. Need to Convert into Intended Shape of three dimensional (32x32x3)

```
# Reshaping Input Array Data

def array_to_pixel_dimensionality_transposition(numarr):
    transposed_list = []

for singlearray in numarr:
    singletransposed = singlearray.reshape(3,32,32).transpose([1, 2, 0])
    transposed_list.append(singletransposed)

return transposed_list
```

- 3. Data Splitting: Already Split between Training and Testing for CIFAR-10 and -100
- 4. Creation of CIFAR-110
 - a. Copy of random sample *of each class of equal size* in CIFAR-10 and append it to CIFAR-100, giving it new classification values
 - Splitting those random samples proportionally to train and testing sets and append them accordingly.

2. b. Data Splitting

- Creation of CIFAR-10 and CIFAR-110

- 1. Data is Already Split between Training and Testing, therefore CIFAR-10 is good to go
- 2. Creation of CIFAR-110
 - a. Copy of random sample *of each class of equal size* in CIFAR-10 and append it to CIFAR-100, giving it new classification values
 - Splitting those random samples proportionally to train and testing sets and append them accordingly.

3. Data Visualization

- Getting a visual of your dataset is always helpful!

```
# sample images
    plt.figure(figsize=(15,15))
    sampleimages = df110train['Image Array'][600:610]
    columns = 5
    for i, image in enumerate(sampleimages):
         plt.subplot(len(sampleimages) / columns + 1, columns, i + 1)
 9
         plt.imshow(image)
10
          10
                              10
                                                   10
          15
                              15
                                                   15
          20
                              20
                                                   20
          25
                              25
                                                   25
                                                                       25
                              30
                 10
                                           20
                                                          10
                                                                              10
                                                                                    20
  20
            0
                      20
                            30
                                                30
                                                               20
                              10
                                                   10
                                                                       10
          15
                              15
                                                   15
```


Final Changes

Low dimensionality (32 x 32 x 3) means difficulty in visualizing images BUT higher speed in terms of machine learning per image.

(fewer pixels means fewer features).

HOWEVER: Data is normalized down (scaling down from (0 to 255) to (0 to 1)

Picture of a Ship:




3. Data Manipulation

- Is there a Preprocessing Mechanism to obtain Sharper images?















Image Rotation









4. Data Reduction Modeling

- Is it Possible to Increase Efficiency via Dimensionality Reduction?
 - Number of features per image = $32 \times 32 \times 3 = 3072$
- Using a dimensionality reduction technique (e.g., PCA) would greatly increase efficiency, but can it be done on data requiring a spatial relationship?
- Using Multi-Layer Perceptron (basic feed forward ANN) for baseline no DR

MLP Accuracy Score on testing set for CIFAR10 is: 43.0%

Time to run MLP model on CIFAR10 is 287 seconds

MLP Accuracy Score on testing set for CIFAR110 is: 1.0%

Time to run MLP model on CIFAR110 is 346 seconds

MLP applied to PCA-Reduced Data

- Determining Amount of Components

```
We want to find the best tradeoff between lost variance and ease of computation,
and so we want to find the lowest number of dimensions that retains variance up to a designated amount
1 1 1
# idea taken from Kagqle user Hamish Dickson for PCA analysis***********
def best PCA fit(fittingdata, minimumretainedvariance=0.857): # recommendation of amount of retained variance
                                                              # to be ~6 out of 7 items
   # model instantiation
   pcacheck = PCA()
    pcacheck.fit(fittingdata)
   cumulativesum = np.cumsum(pcacheck.explained variance ratio )
   num dimensions = np.argmax(cumulativesum >= minimumretainedvariance) + 1 # gotta have at least 1 dimension!
    print("The best number of dimensions to use for PCA on the fitting data is {}.".format(str(num dimensions)))
```


The best number of dimensions to use for PCA on the fitting data is 59. The best number of dimensions to use for PCA on the fitting data is 51.

MLP applied to PCA-Reduced Data

Conclusion:

DR Techniques eviscerate the image pixels' spatial relationship, thus significantly reducing accuracy despite also significantly reducing time.*

*Note: Model convergence was not achieved due to limiting max_iter hyperparameter.



MLP Testing	CIFAR-10 Accuracy	CIFAR-10 Time	CIFAR-110 Accuracy	CIFAR-110 Time
Without PCA	43.0%	287 seconds	1.0%	346 seconds
With PCA	20.0%	32 seconds	1.0%	94 seconds


5. Unsupervised Modeling

- Unsupervised Models used:
 - Spectral Clustering
 - Performs pair-wise similarities between datapoints. This makes it good for non-parametric data such as image data, but pair-wise calculations are always slow.
 - t-Stochastic Neighbor Embedding (t-SNE)
 - Stochastic gradient descent model that is also appropriate for non-parametric data Computationally heavy.
 - <u>Linear Discriminant Analysis (LDA)</u>
 - Creates a hyperplane finding the largest separation between classes (similar to SVM's function in this respect).
 - Requires class input, thus making it not entirely unsupervised.

Spectral Clustering

CIFAR-10 Analysis

CITAIN-10 Allaly	313			
col_0	0	1	2	3
Target Labels				
0	58	80	61	13
1	49	88	58	10
2	41	91	57	15
3	41	78	52	8
4	41	79	63	9
5	31	80	66	7
6	59	87	54	24
7	40	89	61	15
8	34	86	57	19
9	49	70	63	17




Total time to run this model was: 3.4097113609313965 seconds.

Comparison between t-SNE and PCA

<u>Conclusion:</u> Despite the visualization seen earlier, simply graphing a PCA scatter yielded better unsupervised results than t-SNE.

PCA Conversion = ~9.3s t-SNE analysis = 645s

Normally, t-SNE is the model that provides more spread-out results, but there was not full convergence after 300 iterations, taking the above amount of time.



LDA - Setup

LDA with shrinkage done! Time elapsed: 95.24511432647705 seconds LDA without shrinkage done! Time elapsed: 37.49656271934509 seconds LDA on CIFAR-110 with shrinkage done! Time elapsed: 142.13608598709106 seconds Comparison of Accuracies of fitted LDA models - Divided by Number of Features Analyzed Percent Score of Explained Variance of Training Class CIFAR10 w/ shrinkage CIFAR10 w/o shrinkage CIFAR110 w/ shrinkage

Type of Model


LDA Count: 2000 points

Conclusion:

Expected clusters are found:

- 2,3,4(Bird,Cat,Deer)
- 5,7(Dog,Horse)
- 6 alone(Frog)
- 0,8(Airplane, Ship)
- 1,9 (Automobile, Truck)

This shows that there are no errant data trends.



6. Supervised Learning


- Model of Choice: <u>Convolutional Neural Networks (CNNs)</u>

Why? Spatial relationship is taken into account based on its main characteristic: **CONVOLUTIONS**

	23117323113113					
	1,	1,0	1,	0	0	
	O _{×0}	1,	1,0	1	0	4
	0,1	O _{×0}	1,	1	1	
	0	0	1	1	0	
	0	1	1	0	0	
28.	Image			e	Convolved Feature	

Other machine learning models do not use this feature.

Overview of How CNNs Work:



STEPS TO CONVOLUTION

- 1. Define shape of input data
- 2. Define the convolutional planes (called *kernels*)
- 3. Reduce the sample size into a pooling layer (a process called *downsampling*)
- 4. Flatten the downsampled data, place flattened data into dense layers and run model

Overview of CNNs Used:

1. Simple CNN

a. One convolutional layer

2. Multi-Layer CNN

 a. 3 convolutional layers of differing sizes, each with pooling layer immediately after, introduction of "padding"

3. Small Layer CNN

- a. 32-node layers, with 64 dense
- b. Introduction of BatchNormalization

4. Combination CNN

- a. Small layer-multi-CNN
- b. Contains padding and BatchNormalization

A. Simple CNN Architecture

```
# Building Model A.
                                                              Layer (type)
                                                                                             Output Shape
                                                                                                                          Param #
def firstcnn(num classes):
   model = Sequential()
                                                              conv2d 20 (Conv2D)
                                                                                              (None, 30, 30, 32)
                                                                                                                          896
   # First convolutional layer, note the specification of shape
   model.add(Conv2D(32, kernel_size=(3, 3),
                                                              max pooling2d 20 (MaxPooling (None, 15, 15, 32)
                                                                                                                          0
                    activation='relu',
                    input shape=inputshape))
                                                              dropout 23 (Dropout)
                                                                                              (None, 15, 15, 32)
                                                                                                                          0
   model.add(MaxPooling2D(pool size=(2, 2)))
                                                              flatten 12 (Flatten)
                                                                                              (None, 7200)
                                                                                                                          0
   model.add(Dropout(0.25))
   model.add(Flatten())
                                                              dense 23 (Dense)
                                                                                              (None, 128)
                                                                                                                          921728
   model.add(Dense(128, activation='relu'))
   model.add(Dropout(0.25))
                                                              dropout 24 (Dropout)
                                                                                              (None, 128)
                                                                                                                          0
   model.add(Dense(num classes, activation='softmax'))
                                                              dense 24 (Dense)
                                                                                              (None, 10)
                                                                                                                          1290
   model.compile(loss=categorical crossentropy,
                 optimizer=adadelta,
                                                              Total params: 923,914
                                                                                           For CIFAR-10. Total params for
                 metrics=['accuracy'])
                                                              Trainable params: 923,914
                                                                                           CIFAR-110 is 936,814.
   print(model.summarv())
                                                              Non-trainable params: 0
   return model
```

Simple CNN Results


Accuracies for CIFAR-10 (top) and CIFAR-110 (bottom). Highly overfit!

Note: All CNNs trained for 20 epochs

Note 2: Accuracy Graphs do not have uniform y-axis scaling.

Note3: No scaling done to results to place 10-class and 110-class datasets on equal footing.

Simple CNN Results	Accuracy on Testing Set	Total Time to Train
CIFAR-10	64%	806s
CIFAR-110	28%	856s



B. Multi-Layer CNN Architecture

Layer (type)

Output Shape

Param #

```
def multipleconvcnn(num classes):
       model = Sequential()
                                                                                                     conv2d_24 (Conv2D)
                                                                                                                                    (None, 32, 32, 32)
                                                                                                                                                               896
       model.add(Conv2D(32, kernel_size=(3, 3),activation='relu',input_shape=inputshape,padding='same'))
                                                                                                     max pooling2d 24 (MaxPooling (None, 16, 16, 32)
                                                                                                                                                               0
       model.add(MaxPooling2D((2, 2),padding='same'))
       model.add(Dropout(0.25))
                                                                                                     dropout 31 (Dropout)
                                                                                                                                    (None, 16, 16, 32)
                                                                                                                                                               0
       model.add(Conv2D(64, (3, 3), activation='relu',padding='same'))
                                                                                                     conv2d 25 (Conv2D)
                                                                                                                                    (None, 16, 16, 64)
                                                                                                                                                               18496
       model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
       model.add(Dropout(0.25))
                                                                                                     max pooling2d 25 (MaxPooling (None, 8, 8, 64)
                                                                                                                                                               0
       model.add(Conv2D(128, (3, 3), activation='relu',padding='same'))
                                                                                                     dropout 32 (Dropout)
                                                                                                                                    (None, 8, 8, 64)
                                                                                                                                                               0
       model.add(MaxPooling2D(pool size=(2, 2),padding='same'))
10
                                                                                                     conv2d 26 (Conv2D)
       model.add(Dropout(0.25))
                                                                                                                                    (None, 8, 8, 128)
11
                                                                                                                                                               73856
12
       model.add(Flatten())
                                                                                                     max pooling2d 26 (MaxPooling (None, 4, 4, 128)
                                                                                                                                                               0
13
       model.add(Dense(128, activation='relu'))
                                                                                                     dropout 33 (Dropout)
                                                                                                                                    (None, 4, 4, 128)
14
       model.add(Dropout(0.25))
                                                                                                                                                               0
15
       model.add(Dense(num classes, activation='softmax'))
                                                                                                     flatten 16 (Flatten)
                                                                                                                                    (None, 2048)
                                                                                                                                                               0
16
                                                                                                     dense 31 (Dense)
                                                                                                                                    (None, 128)
                                                                                                                                                               262272
17
       model.compile(loss=categorical crossentropy,
18
                     optimizer=adadelta,
                                                                                                     dropout 34 (Dropout)
                                                                                                                                    (None, 128)
                                                                                                                                                               0
19
                     metrics=['accuracy'])
                                                                                                     dense 32 (Dense)
                                                                                                                                    (None, 10)
                                                                                                                                                               1290
20
                                                           Total params for
21
       print(model.summary())
                                                                                                     Total params: 356,810
                                                           CIFAR-110 is 369,710.
                                                                                                     Trainable params: 356,810
22
                                                           (less than Simple CNN).
                                                                                                     Non-trainable params: 0
23
       return model
```

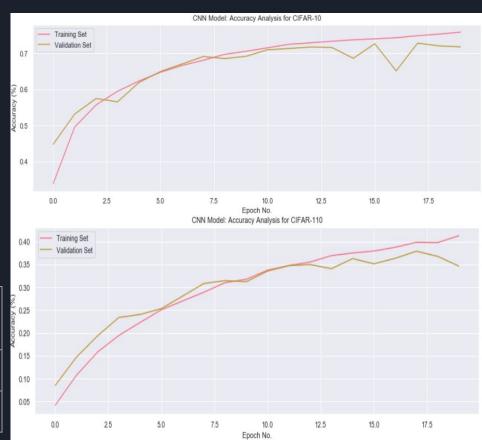
Multi-Layer CNN Results

<u>padding='same'</u> means that part of the kernel convolutes over the outside of the image for a uniformly-shaped convolution (essentially adding another row and column of pixels, each with "0".

Much less overfit than Simple CNN (although still slightly overfit).

Significantly more time to run than Simple CNN!

Multi-Layer CNN Results	Accuracy on Testing Set	Total Time to Train
CIFAR-10	72%	1170s
CIFAR-110	35%	1513s



C. Small-Layer CNN Architecture

```
def small layer cnn(num classes):
                                                                           Layer (type)
                                                                                                     Output Shape
                                                                                                                              Param #
    model = Sequential()
                                                                           conv2d 43 (Conv2D)
                                                                                                      (None, 30, 30, 32)
                                                                                                                              896
    model.add(Conv2D(32, kernel size=(3, 3),
                                                                           batch normalization 5 (Batch (None, 30, 30, 32)
                                                                                                                             128
                      activation='relu'.
                      input shape=inputshape))
                                                                           max pooling2d 43 (MaxPooling (None, 15, 15, 32)
                                                                                                                              0
    model.add(BatchNormalization())
                                                                           dropout 59 (Dropout)
                                                                                                      (None, 15, 15, 32)
                                                                                                                              0
    model.add(MaxPooling2D(pool size=(2, 2)))
    model.add(Dropout(0.3))
                                                                           conv2d 44 (Conv2D)
                                                                                                      (None, 15, 15, 32)
                                                                                                                              9248
    model.add(Conv2D(32, (3, 3), activation='relu',padding='same'))
                                                                                                      (None, 15, 15, 32)
                                                                           conv2d 45 (Conv2D)
                                                                                                                              9248
    model.add(Conv2D(32, (3, 3), activation='relu',padding='same'))
    model.add(BatchNormalization())
                                                                           batch normalization 6 (Batch (None, 15, 15, 32)
                                                                                                                             128
    model.add(MaxPooling2D(pool size=(2, 2)))
                                                                           max pooling2d 44 (MaxPooling (None, 7, 7, 32)
    model.add(Dropout(0.3))
                                                                                                                              0
    model.add(Flatten())
                                                                           dropout 60 (Dropout)
                                                                                                      (None, 7, 7, 32)
                                                                                                                              0
    model.add(Dense(64, activation='relu'))
    model.add(Dropout(0.4))
                                                                           flatten 25 (Flatten)
                                                                                                      (None, 1568)
                                                                                                                              0
    model.add(Dense(num classes, activation='softmax'))
                                                                           dense 49 (Dense)
                                                                                                      (None, 64)
                                                                                                                              100416
    model.compile(loss=categorical crossentropy,
                                                                           dropout 61 (Dropout)
                                                                                                      (None, 64)
                                                                                                                              0
                   optimizer='rmsprop',
                                                                           dense 50 (Dense)
                                                                                                      (None, 10)
                                                                                                                              650
                   metrics=['accuracy'])
    print(model.summary())
                                                                                                     Total params for CIFAR-110
                                                                           Total params: 120,714
```

Trainable params: 120,586

ic 127 71/

return model

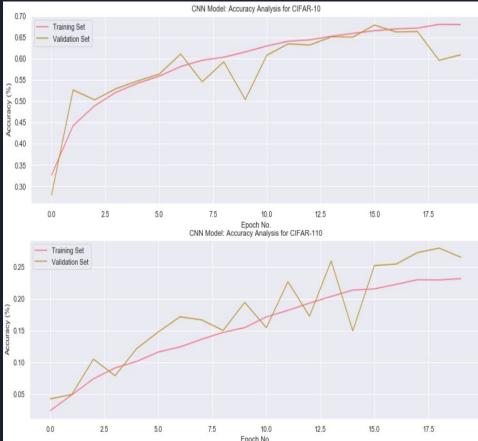
Small-Layer CNN Results

<u>BatchNormalization():</u> As the name implies, it initializes the parameters with zero mean and unit variance.

Sporadic validation accuracies are seen here. Perhaps the holdout (of 10% of the training set) is too small for consistent results, although this was not a problem for the other CNNs.

Training times increased to 40 and 50 minutes (using 100% CPU and several GB of RAM).

Small-Layer CNN Results	Accuracy on Testing Set	Total Time to Train
CIFAR-10	61.05%	2434s
CIFAR-110	26.22%	3102s



D. "Combination" CNN Architecture

```
combo cnn(num classes):
                                                                                     Layer (type)
                                                                                                                     Output Shape
model = Sequential()
                                                                                     -----
                                                                                     conv2d 55 (Conv2D)
                                                                                                                     (None, 32, 32, 32)
                                                                                                                                                  1568
model.add(Conv2D(32, kernel size=(4, 4),activation='relu',input shape=inputshape.padding='same')
model.add(BatchNormalization())
                                                                                     batch normalization 15 (Batc (None, 32, 32, 32)
                                                                                                                                                  128
model.add(MaxPooling2D((2, 2),padding='same'))
                                                                                     max pooling2d 53 (MaxPooling (None, 16, 16, 32)
                                                                                                                                                  0
model.add(Dropout(0.3))
model.add(Conv2D(32, (3, 3), activation='relu',padding='same'))
                                                                                     dropout 73 (Dropout)
                                                                                                                     (None, 16, 16, 32)
                                                                                                                                                  0
model.add(BatchNormalization())
                                                                                     conv2d_56 (Conv2D)
                                                                                                                     (None, 16, 16, 32)
                                                                                                                                                  9248
model.add(MaxPooling2D(pool size=(2, 2),padding='same'))
model.add(Dropout(0.3))
                                                                                     batch normalization 16 (Batc (None, 16, 16, 32)
                                                                                                                                                  128
model.add(Conv2D(64, (3, 3), activation='relu',padding='same'))
                                                                                     max_pooling2d_54 (MaxPooling (None, 8, 8, 32)
                                                                                                                                                  0
model.add(BatchNormalization())
                                                                                     dropout 74 (Dropout)
model.add(MaxPooling2D(pool size=(2, 2),padding='same'))
                                                                                                                     (None, 8, 8, 32)
                                                                                                                                                  0
model.add(Dropout(0.4))
                                                                                     conv2d 57 (Conv2D)
                                                                                                                     (None, 8, 8, 64)
                                                                                                                                                  18496
                                                                                     batch_normalization_17 (Batc (None, 8, 8, 64)
                                                                                                                                                  256
model.add(Flatten())
model.add(Dense(128, activation='relu'))
                                                                                     max pooling2d 55 (MaxPooling (None, 4, 4, 64)
                                                                                                                                                  0
model.add(Dropout(0.5))
                                                                                     dropout 75 (Dropout)
                                                                                                                     (None, 4, 4, 64)
                                                                                                                                                  0
model.add(Dense(num classes, activation='softmax'))
                                                                                    flatten 29 (Flatten)
                                                                                                                     (None, 1024)
                                                                                                                                                  0
model.compile(loss=categorical crossentropy,
                                                                                     dense 57 (Dense)
                                                                                                                     (None, 128)
                                                                                                                                                  131200
            optimizer='adam',
            metrics=['accuracy'])
                                                                                     dropout 76 (Dropout)
                                                                                                                     (None, 128)
                                                                                                                                                  0
                                                                                     dense 58 (Dense)
                                                                                                                     (None, 10)
                                                                                                                                                  1290
print(model.summary())
                                                                                     _____
                                                                                                                    Total params for CIFAR-110
                                                                                     Total params: 162,314
                                                                                     Trainable params: 162,058
                                                                                                                   is 175.214.
return model
                                                                                     Non-trainable params: 256
```

Combination CNN Results

Applied BatchNormalization, padding, and multiple small layers (i.e., the combination of Models A, B, & C).


Least overfit of all the models, with CIFAR-110 training being *underfit*.


Despite normalizations meant to streamline/speed up training, this model took the longest time to train!

Combination CNN Results	Accuracy on Testing Set	Total Time to Train
CIFAR-10	59.55%	3032s
CIFAR-110	26.04%	3396s




7. CONCLUSION - Accuracies






7. CONCLUSION - Training Times



7. CONCLUSION - Efficiency Scores (measured in Accuracy / Min. Rates)



7. CONCLUSION - FINAL ANALYSIS

<u>The Problem:</u> Which supervised model proves to be the best accuracy/efficiency tradeoff for image classification?

The Answer: Out of the four CNN models tested, the

Multi-Layer CNN Keras Model

is the best to use for both small classification sets (CIFAR-10) and large classification sets.

Explanation: The Simple CNN, as the name would suggest, converged the fastest. This gave it the *best efficiency score out of all models* for both CIFAR-10 and CIFAR-110 datasets. At the same time, the simple CNN **was also the most overfit of the four models**. The Multi-Layer CNN was the least overfit when applied to both CIFAR-10 and CIFAR-110 datasets. It also was the second fastest and thus had the second highest efficiency score.

F.A.Q.

- 1. Why weren't other metrics used, including precision, specificity, or F1 scoring?
 - These metrics make more sense for a binary classifier: what exactly is a "false negative" for a multi-classification problem??? For this reason, accuracy was focused on (and its modelling analog, loss function calculation).
- 2. Why is training the focus here? Once a model is trained and sample weights are obtained for a given model, wouldn't making testing predictions be the point of focus for speed?
 - Maybe, but focusing on training time nonetheless speaks to this anyway. There is no good way to measure prediction time, since it happens so fast. The fractions of a second to make a prediction should be proportionate (not equal) to training time. If the focus of the project is efficiency, then having a much larger scale to analyze, i.e., training time, would be more appropriate than testing time.
- 3. What was the purpose of combining CIFAR-10 and CIFAR-100 to create CIFAR-110?
 - Much like MNIST and other popular image datasets where the data can literally be called by a keras import - the CIFAR datasets have been analyzed ad nauseum. I wanted to explore how the introduction of new classes would affect CIFAR-110, especially when those new classes are relatively similar to some of those 100 classes. In short, this is an added challenge for modeling purposes to buttress the results of those models.

Thank You for your attention.

Any questions?

(this is the part where you run)



